

AI Foundation

Knowledge Graph

Introduction to Knowledge Graph and its
applications

Sept 30, 2025

LESSON PLAN

01 Introduction to Knowledge Graph

- Components

02 GraphRAG vs. RAG


- Limitation on RAG
- Main differences: GraphRAG vs. RAG

03 GraphRAG - Under the hood

04 How to build a Knowledge Graph

Define Knowledge Graph

Example



(Person) - [Teaches] →

(Session)

← [Attends] -

(Person)

Component #1: Node

**Node
(Entity)**

- a data record
- represents entities (people, thing, concept, etc)
- has label, e.g. **Person**
- key-value property:
 - **name "Chrissy"**

Chrissy


Memory

Component #2: Edge

Edge (Relationship)

- represents relationship
- connection b/w two nodes
- **has one direction**
- source vs. target node


- **Taxonomic:** `:IS_A`, `:SUBCLASS_OF` (usually transitive)
- **Part-whole:** `:PART_OF`, `:HAS_PART` (often transitive)
- **Membership:** `:MEMBER_OF`, `:BELONGS_TO`
- **Attribution / property-as-edge:** `:HAS_ATTRIBUTE`, `:HAS_STATUS` (or use `DatatypeProperty` in RDF)
- **Identity / equivalence:** `:SAME_AS` (e.g., `owl:sameAs`)
- **Causal / dependency:** `:CAUSES`, `:REQUIRES`, `:DEPENDS_ON`
- **Temporal:** `:VALID_DURING`, `:HAPPENED_BEFORE/AFTER` (or put timestamps on the edge)
- **Spatial:** `:LOCATED_IN`, `:NEAR`
- **Event roles** (n-ary relations via an event node): `(:Event)<--[:AGENT]-(:Actor)`, `(:Event)-[:PATIENT]->(:Asset)`
- **Provenance:** `:CREATED_BY`, `:DERIVED_FROM`, `:CITED_BY`
- **Similarity / relatedness:** `:SIMILAR_TO` (often with a `score` / `weight`)



Component #3: Properties

Properties

- additional info about nodes & edges
- Data types: string, integer, float, array, boolean, etc



Knowledge graph is..

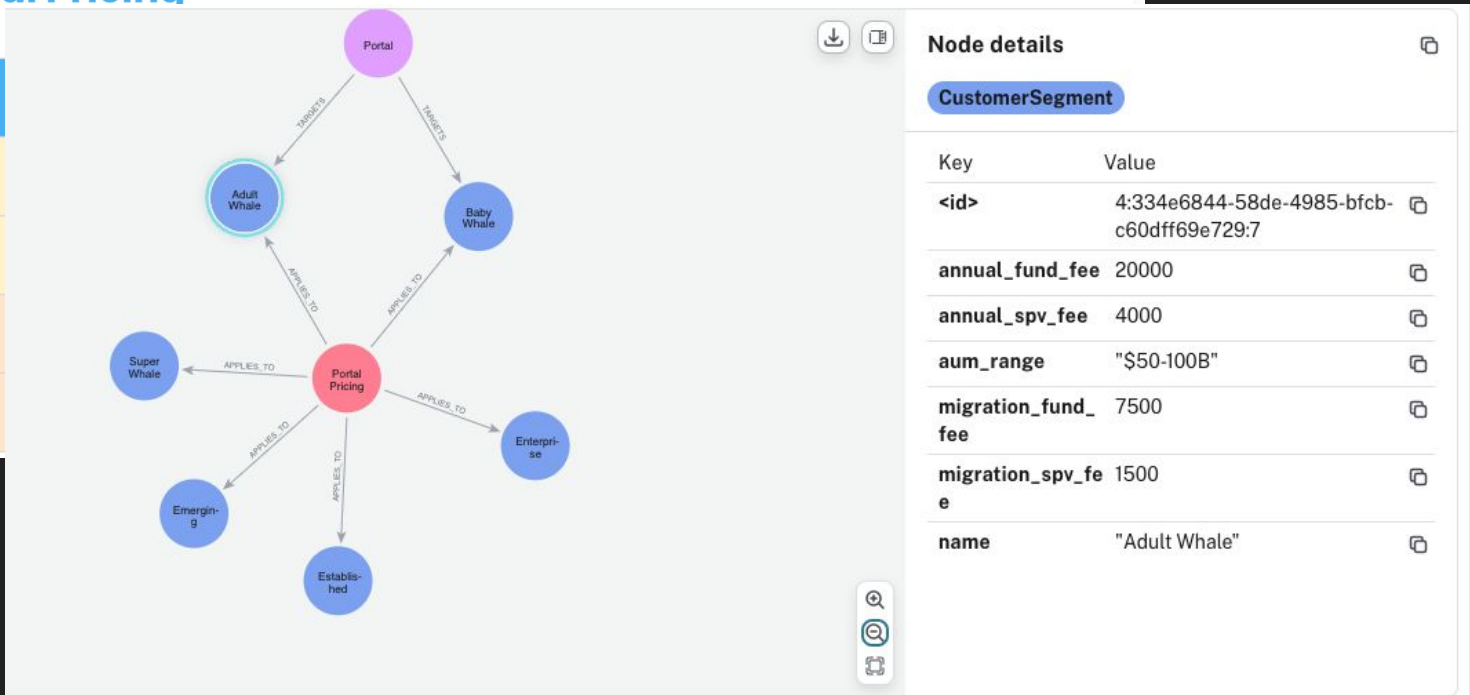
- A database that stores information in **nodes** and **edges**.
 - Both nodes and edges have **properties**.
 - Edges have **direction** and **types**.

(Source: Deep Learning AI - Knowledge Graph for RAG)

Example on Neo4j

Proposed Portal Pricing

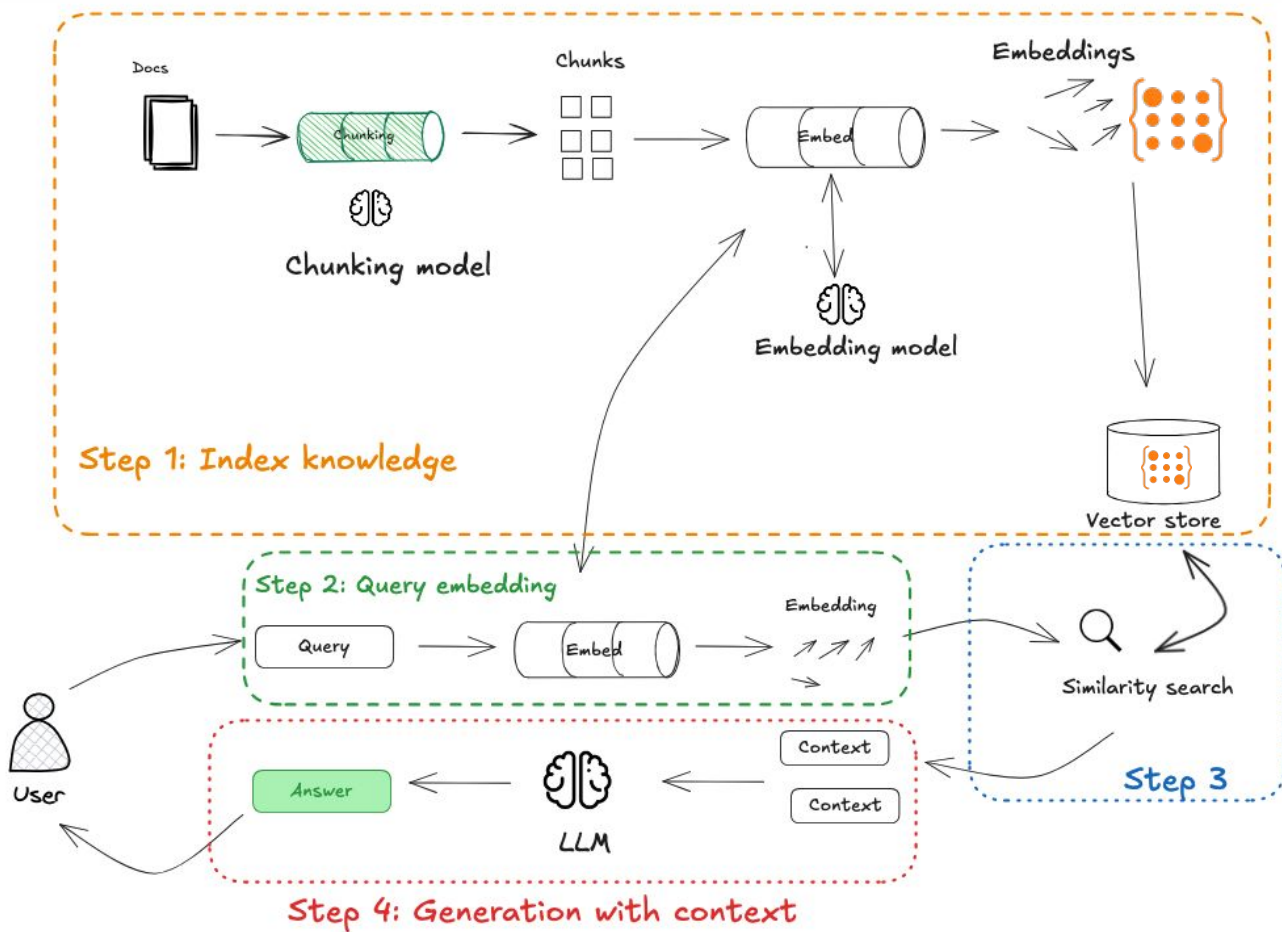
Fee Type	Unit
Annual Fee (Recurring)	Per-Fund
	Per-SPV
Data Migration Fee (One-time)	Per-Fund
	Per-SPV



Deep dive into Knowledge Graph setup

The idea of GraphRAG and more

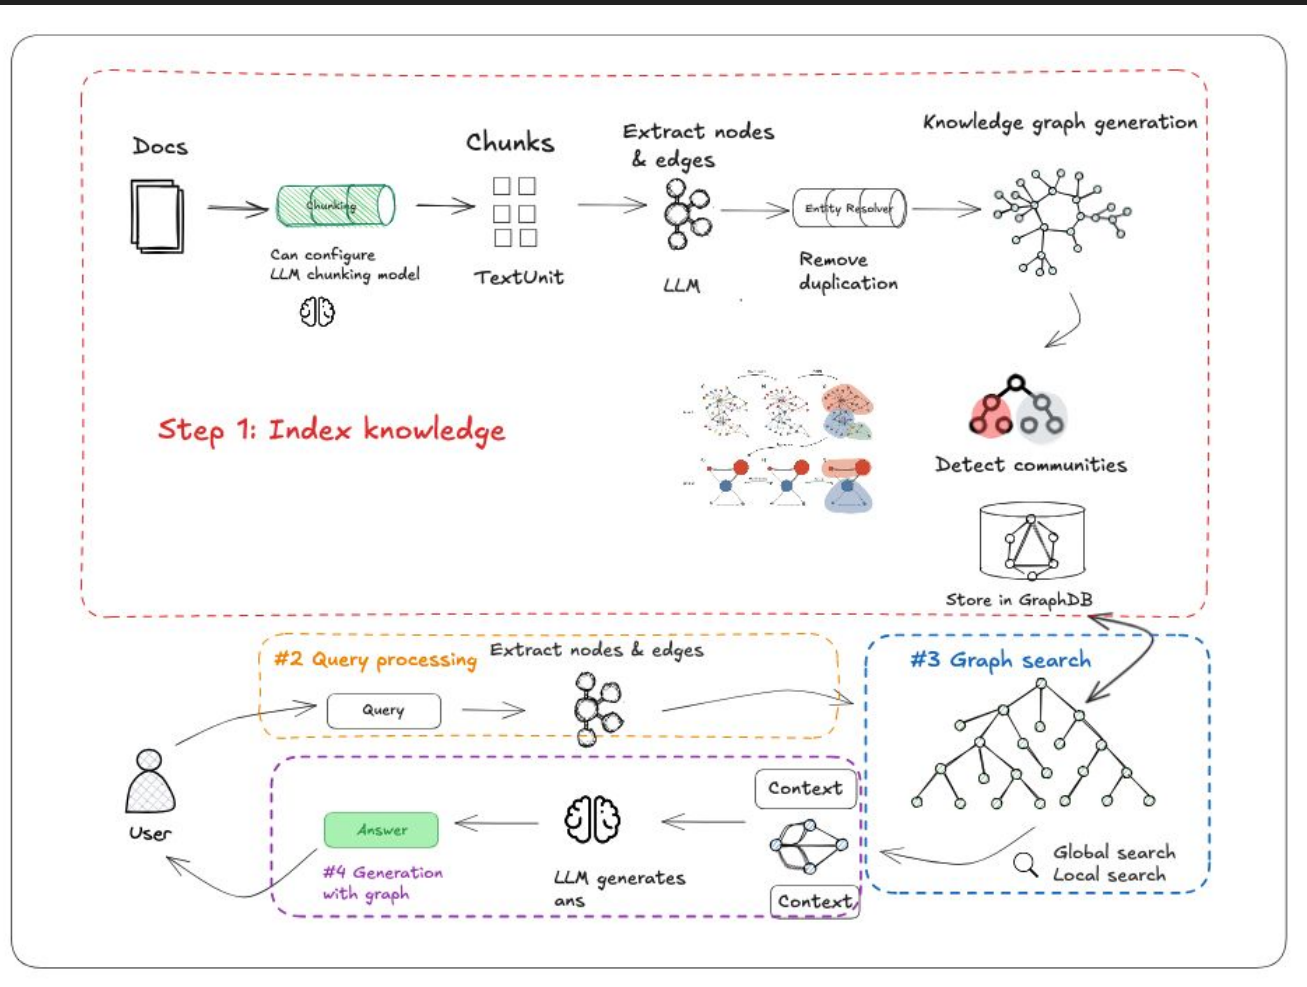
RAG



When RAG fails to deliver

- Simple keyword/semantic search may miss complex relationships
- Difficulty with multi-hop reasoning
- Limited context understanding across document boundaries
- Challenges with entity disambiguation

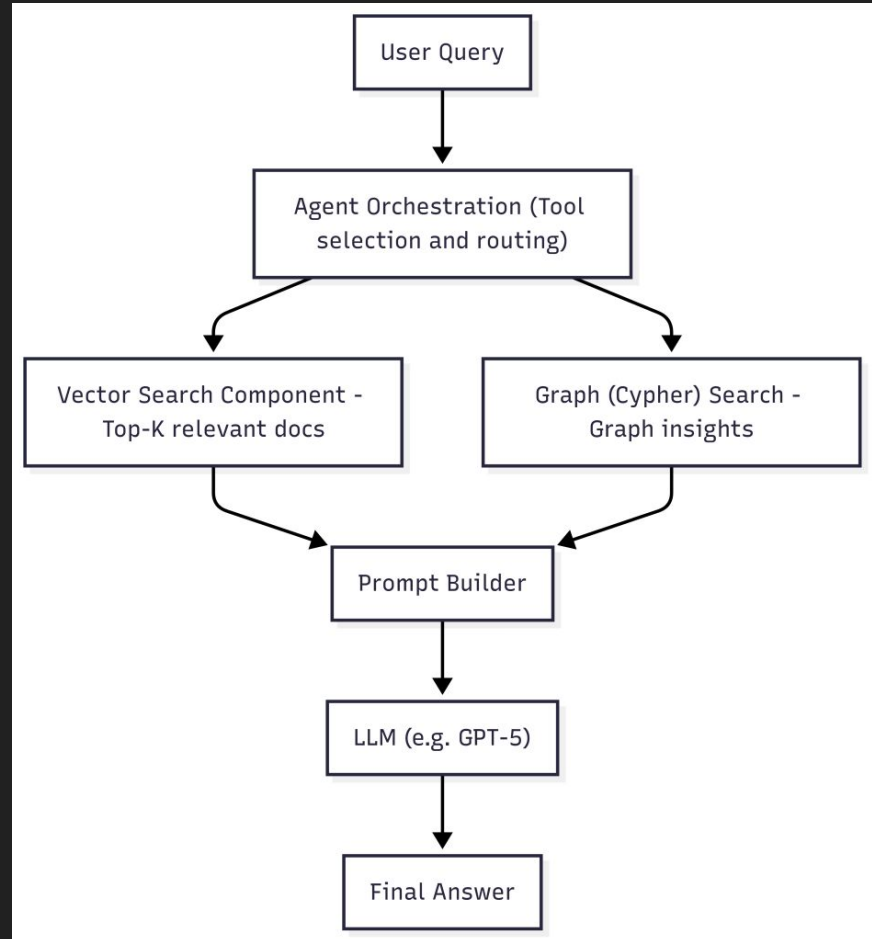
Here comes Graph RAG



GraphRAG architecture

This combination of a knowledge graph and vector search is particularly effective when your application requires:

- Reasoning over complex architectures (e.g., microservices, IT assets, workflows)
- Querying both structured metadata and unstructured documentation
- Combining multiple sources into one coherent system



Compare RAG vs. GraphRAG

Aspect	Traditional RAG	GraphRAG
Structure	Flat vector search	Graph-based relationships
Context	Individual chunks	Connected knowledge network
Reasoning	Similarity-based	Relationship-based
Multi-hop	Limited	Excellent
Complexity	Lower	Higher

GraphRAG Glossary

Text Unit - a chunk of text. The size of these chunks, their overlap can be configured.

NER - Named-Entity-Recognition

Communities - a group of nodes that are more densely connected to each other than to nodes outside the group.

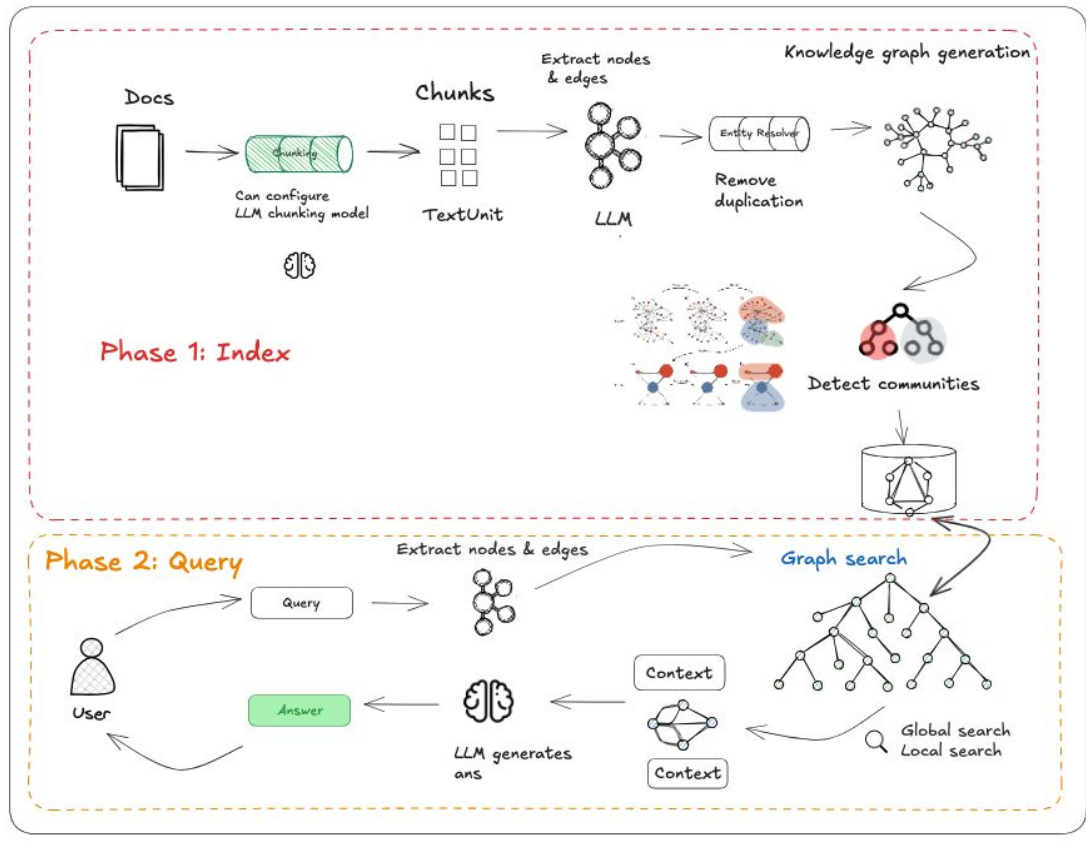
Community Report - The contents of each community are summarized into a generated report, useful for human reading and downstream search.

Community level - the hierarchy tier of clustering; level 0 is the most generic grouping, and higher numbers mean finer, more specific sub-communities.

Leiden - algorithm used to perform community detection.

GraphRAG – Under the hood

What's included in a GraphRAG?



Phase #1: Indexing – Graph Generation

In this phase, there are a few steps:

1. Document → **Chunks**
2. LLM **extracts** entities + relationships from each chunk
3. Merge overlapping entities across chunks (entity resolution)
4. Build knowledge graph from consolidated entities/relationships
5. Apply Leiden algorithm → Detect hierarchical communities
6. Generate community summaries at multiple levels (bottom-up)
7. Create community reports describing each cluster's themes

#1 Text chunking



There are multiple chunking strategies:

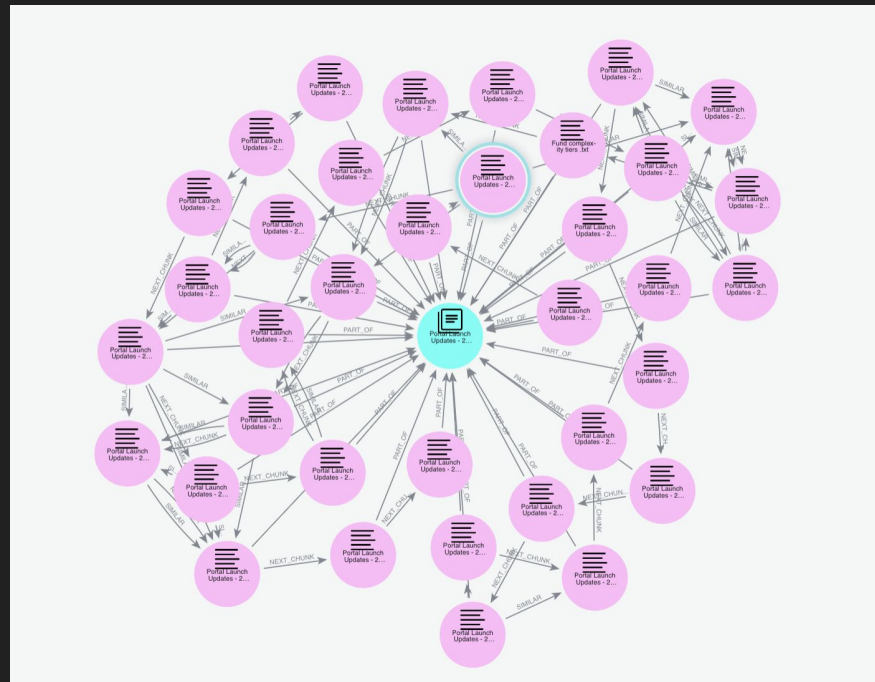
- **Splitting:** (common)

- Split a document into equal sized sections (by character or token-count),
- with an optional overlap (typical sizes are 250-500 tokens with 50-100 tokens overlap)

- **Hierarchical Document Chunking:**

- Split a document alongside lexical boundaries - chapters, sections, paragraphs

- **Sentence Chunking:** Split a document into individual sentences



Why chunking?

Why chunking?

- Still need to chunk so LLM does not need to handle all at once

Is chunking same as RAG?

- Same chunking but embedding stored in vector database (optional)


What is the output after chunking?

- Text Units then go through entity extraction process.

#2 Graph extracting



- After chunking, next we will extract nodes and edges using LLM
- Apply Named-entity recognition (NER) to recognize entities and remove duplication
- The default graph shape is lexical graph*




Source: Neo4j

#3 Community detection, then generate community reports (optional)

- Apply Hierarchical Leiden Algorithm to detect communities
- Generate summaries -> report

([source](#))



#4 Store in graph database


Database information

Nodes (430)

* **__Entity__** Account AUMRange Campaign
Characteristic Chunk Company Competitor
CustomerGroup CustomerSegment DiscountStructure
Document DocumentType EntityType Event
Feature Fee FeeType FinancialEntity
[Show all \(26 more\)](#)

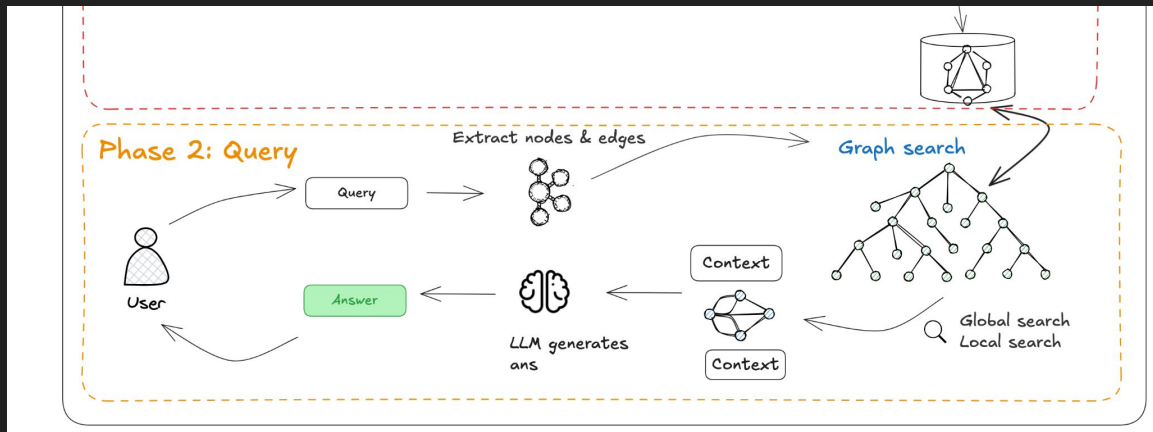
Relationships (1,206)

* ADD_ON_TO ALSO_KNOWN_AS APPLIES_TO
APPLIES_TO_ENTITY ASSOCIATED_WITH
belong_to_subcategory BELONGS_TO BETTER_THAN
BUNDLED_WITH CAN_BUNDLE_WITH
COMMUNICATES_TO COMPETES_WITH
CONSISTS_OF CONSULTED_BY CONTAINS
COORDINATES_ONBOARDING_FOR
CREATES_LANDING_PAGE_FOR
CREATES_PRESS_RELEASES_FOR DESCRIBES



Phase #2: Query


1. Extract query entities (using LLM)
2. Perform graph retrieval techniques:
 - a. Entity-based
 - b. Community-based
 - c. Text-to-Cypher
 - d. Others
3. Generate final response using retrieved context



More techniques can be viewed [here](#)

For entity-based retrieval

Graph traversal algorithms
(methods like **breadth-first search aka BFS**):
explores a graph by visiting all the neighboring nodes of a given level before moving to the next level




Text-to-Cypher

Query: What is the portal pricing for enterprise segment for funds?

```
1 MATCH (portal:Product {name: 'Portal'})-[:TARGETS]->(segment:CustomerSegment {name: 'Enterprise'})
2 RETURN
3   portal.name AS Product,
4   segment.name AS Segment,
5   segment.aum_range AS AUM_Range,
6   segment.annual_fund_fee AS Annual_Fund_Fee,
7   segment.migration_fund_fee AS Migration_Fund_Fee,
8   (segment.annual_fund_fee + segment.migration_fund_fee) AS Total_First_Year_Cost
```

	Product	Segment	AUM_Range	Annual_Fund_Fee	Migration_Fund_F	Total_First_Year
1	"Portal"	"Enterprise"	"\$1B-15B"	15000	7500	22500

Output after graph search



Node details

CustomerSegment

Key	Value	
<id>	4:334e6844-58de-4985-bfcb-c60dff69e729:5	
annual_fund_fee	15000	
annual_spv_fee	3000	
aum_range	"\$1B-15B"	
migration_fund_fee	7500	
migration_spv_fee	1500	
name	"Enterprise"	

Community-based retrieval - global search (source)

Global search with dynamic community selection

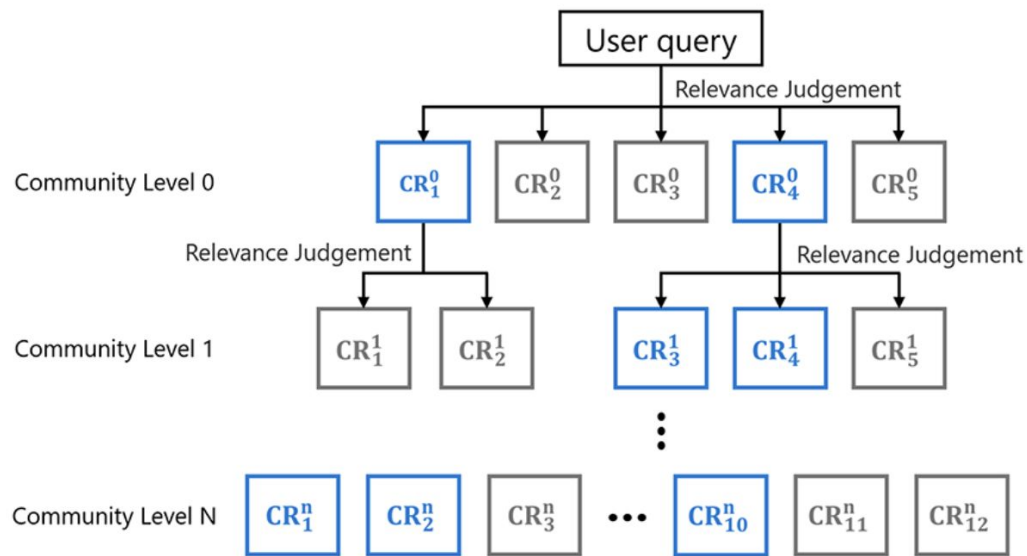


Figure 1: Dynamic community selection workflow

Step 1: Recursive Community Selection

Start at root level (level 0), using an LLM to rate the relevance of each community report (CR) against the user query.

If the community report is not relevant (grey boxes), then we move on to the next report.

If the community report is relevant, then we repeat the same operation to its child nodes. We keep the report (blue boxes) if none of its child communities is relevant or if it is a leaf node.

Step 2: Map-reduce Search

Perform map-reduce search on the selected community reports to generate the final response.

GraphRAG Glossary - p.2

Global search - handle queries that span the entire dataset

Local search - optimize for targeted queries , drawing from smaller subset of documents that closely match the user's input

DRIFT search - a hybrid b/w global and local search

Breadth-first search (BFS): explores a graph by visiting all the neighboring nodes of a given level before moving to the next level

Cypher - a graph-based query language developed by Neo4j

Example on Neo4j



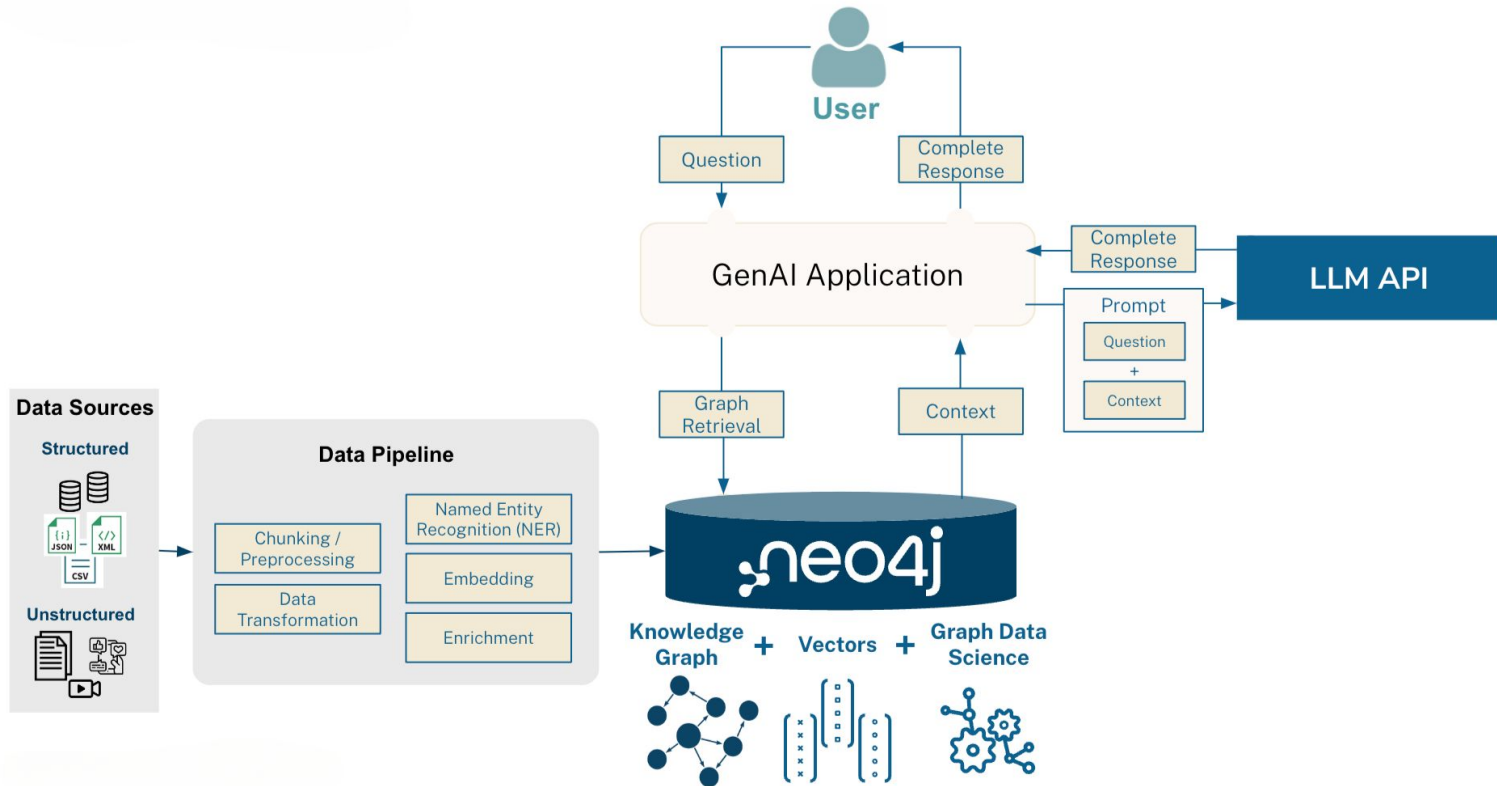
Direct Answer: The Portal pricing for the **Enterprise** segment is composed of an **Annual Fund Fee** of \$15,000 and a **Migration Fund Fee** of \$7,500, resulting in a **Total Pricing** of \$22,500.

Key Insights:

- The total cost for the **Enterprise** segment stands at **\$22,500** annually.
- The **Annual Fund Fee** constitutes approximately **66.67%** of the total pricing, while the **Migration Fund Fee** accounts for **33.33%**.



Graph RAG Neo4j



Part 3: How to build Knowledge Graph

Use cases & implementations

How-to

Step 1: Prepare data: structured and unstructured

Step 2: Configure the following items:

- chunking technique: sentence, semantic or else
- embedding model
- LLM used to generate graph: LlamaIndex, Langchain

Step 3: Determine the database

- DB for vector
- DB for graph: GraphRAG, Neo4j, Memgraph, LightRAG
- Hybrid: Neo4j

Demo workflow (Flowise)

[Link](#)

Demo workflow (ZEP)

- Store long-term memory as graph and retrieve for responses
- N8n workflow using ZEP: [link](#)
- View graph on ZEP: [Link](#)

KG Application

From legal documents to knowledge graph:[link](#)

KG Terminology

For reference

Term	Simple Definition
Entity	A thing or concept, e.g. a person, company, location or product
Node	A dot in the graph - usually represents an entity
Edge	A line in the graph - represents the relationship between two entities
Triple	A statement in the format: Subject -> Predicate -> Object
Subject	The entity doing or being described
Predicate	The relationship or property
Object	The target of the relationship

Links & Resources

- Knowledge Graphs for RAG - DeepLearning.AI ([link](#))
- From Local to Global ([link](#))
- GraphRAG ([link](#))
- Demo workflow for KG ([link](#))
- Agentic RAG ([link](#))

Product Knowledge Library

Structure KG for PKL on Neo4j DB

- document type
- document link



belongs_to



belongs_to



is_about

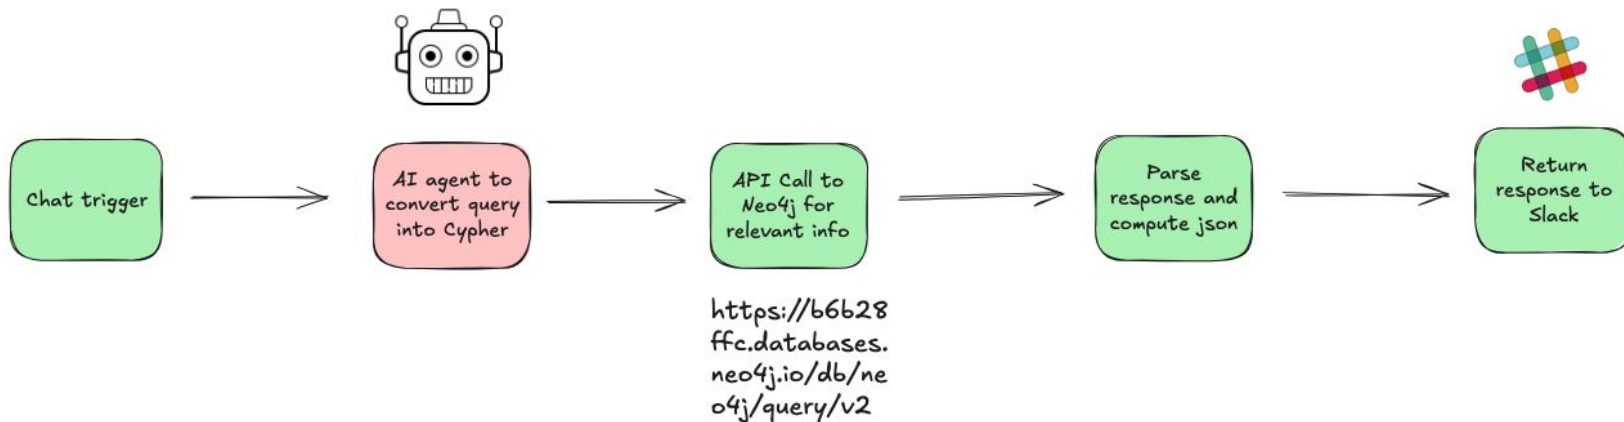
Product name

- All
- Integration
- Side Letter
- E-signature
- Data Room
- IDM
- Investor Portal
- Platform
- Investor Access
- AAA
- Fundsub
- OCR
- Landing Page

sub-product name

Diagram

n8n sub-workflow to showcase KG



For reference

Part : Graph solutions



Neo4j, Memgraph, GraphRAG (Microsoft) &
LightRAG

The GraphRAG ecosystem

neo4j



GraphRAG



LightRAG